

The Relational Problem

Situation:

- Big Data is exploding all around us
- The explosion of Big Data is growing at an increasing rate

Problem:

Relational Database Systems cannot handle the volume, velocity and variety of Big Data.

The Relational Problem

Challenges

- How can we collect, process and store so much data?
- How can we wade through so much data in a timely manner?
- How can we analyze so much data and gain meaningful insights?
- Can my existing systems/architectures handle this?
 - Why not?
- Can my existing staff (skills & tools) make the transition?

The Relational Problem

The Relational Problem

- The Relational Database: 40+ year-old technology
 - Demands STRUCTURE: Tables, Rows, Columns, Keys, Indexes
 - Demands ACID Transaction Compliance
 - Keep my data consistent across transactions
 - Consistency VS Fast Performance and Throughput
 - RDBMS Software Choices
- Traditional Database Server Architecture
 - Memory, CPU, Storage
 - The cost of scaling “UP”

The Relational Problem

Techniques to Help Handle Big Data in RDBMS

- Scaling – Up versus Out
 - UP: Add CPU, Memory, Storage
 - OUT: Add server nodes to a cluster
 - “Linear” scalability
- Clustering
 - “Commodity” Hardware
 - Introduces some overhead for inter-node communications
- Replication
- Parallelization
- Sharding
- Move the COMPUTE closer to the DATA

Side Notes -- Scaling Out

NOTE:

Be sure to watch the Google Container Data Center video. (2009, getting started)

<https://www.youtube.com/watch?v=zRwPSFpLX8I>

Challenges:

- Heat
- Maintenance
- UPS – Uninterruptible Power Supply

Side Notes -- Scaling Out

NOTE:

**Be sure to also watch the Google Data Center tours
(9 years later)**

<https://www.youtube.com/watch?v=zDAYZU4A3w0>

<https://www.youtube.com/watch?v=XZmGGAbHqa0>

<https://www.youtube.com/watch?v=avP5d16wEp0>

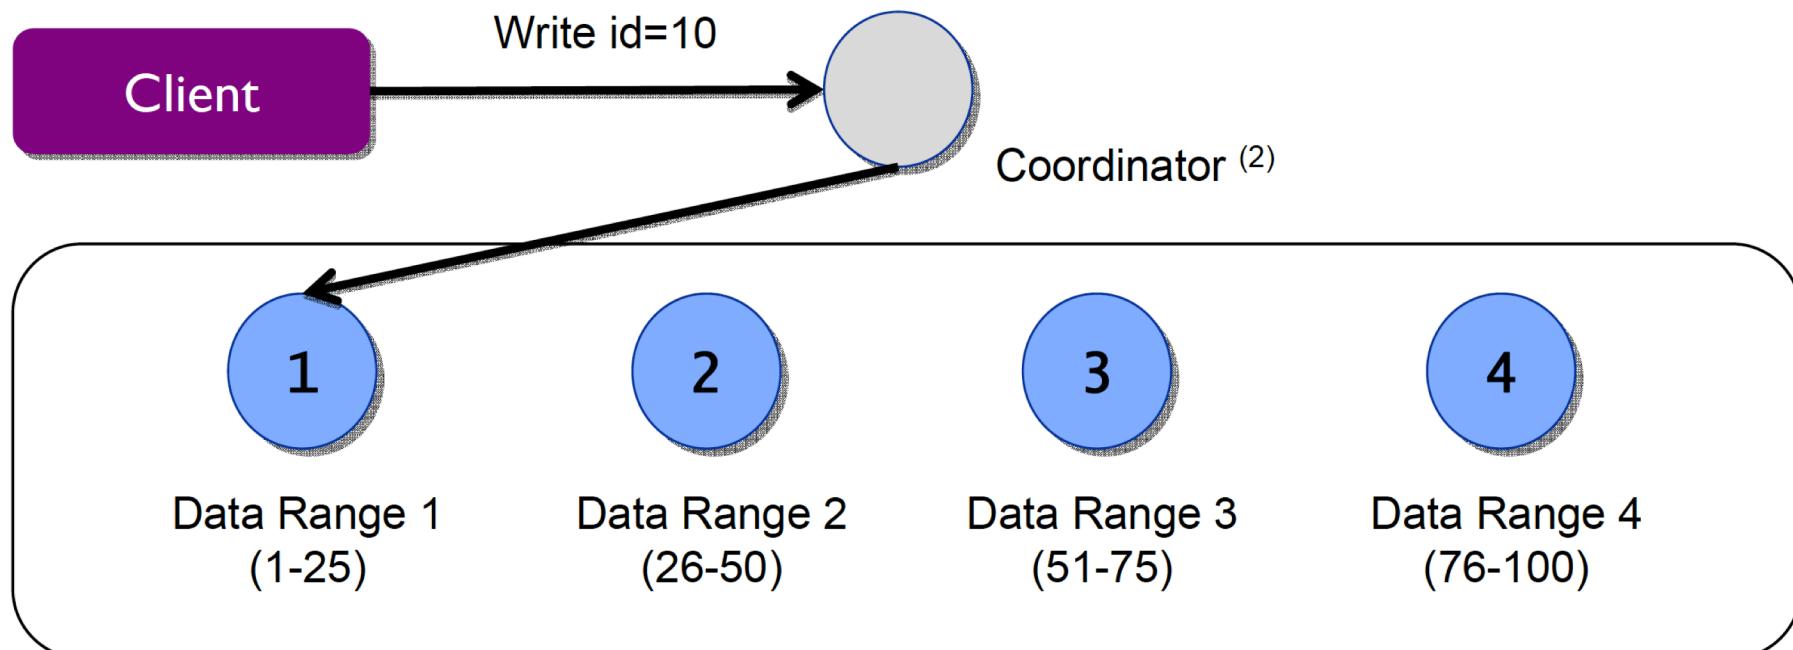
The Relational Problem

Techniques to Help Handle Big Data in RDBMS

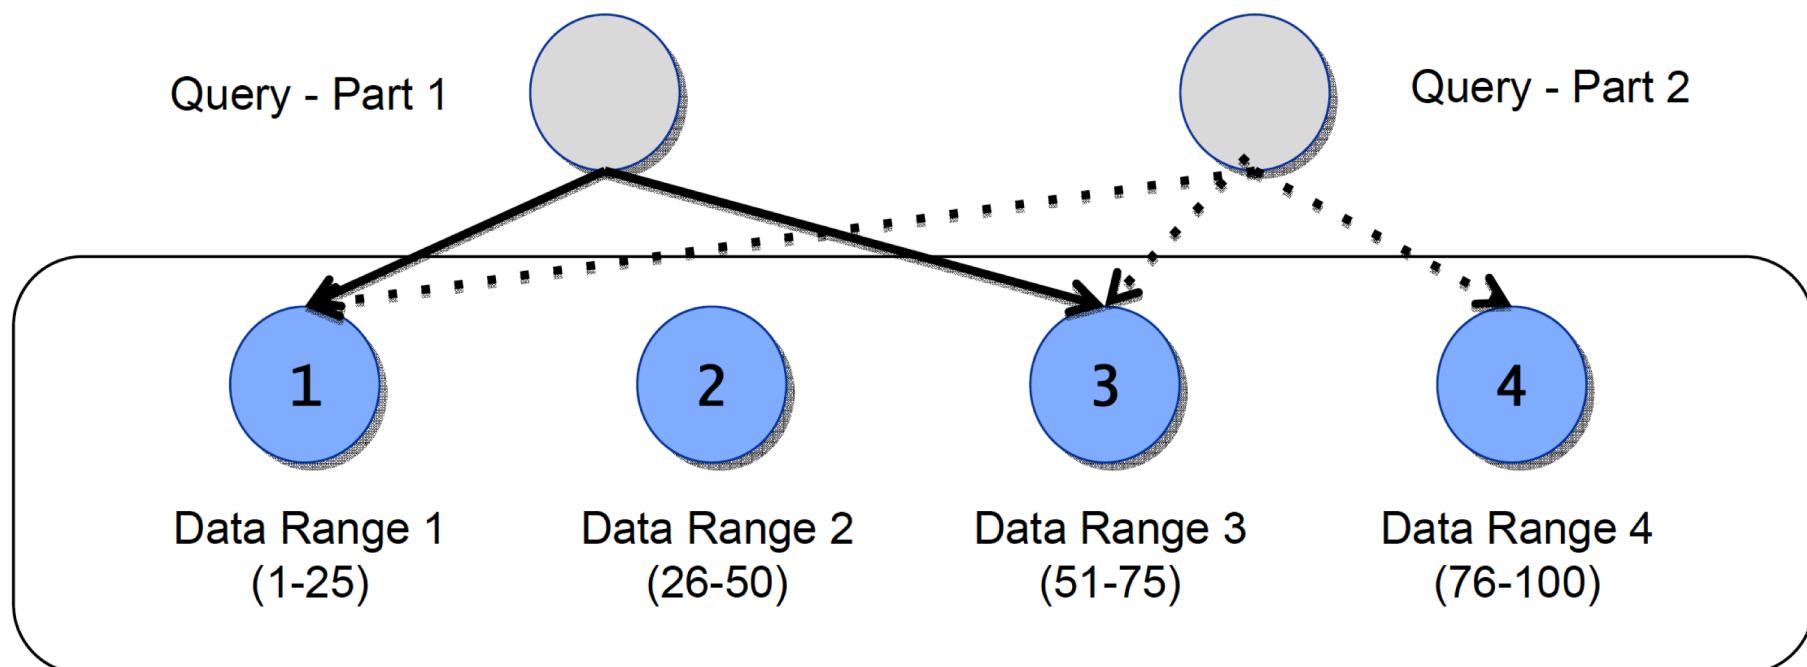
- Clustering without sharding for
 - Built Upon Simple Replication
 - HA - High Availability
 - No SPOF – Single Point of Failure
- Master-Master, Master-Slave Replication
 - The need for a “Watcher”
 - Problems with Geographically Distributed clusters

- ◆ **Scale Out**: Use large clusters of commodity hardware
 - As opposed to **scale up** (with heavy duty servers)
 - Often done in the cloud for ease of startup and growth
- ◆ **Distribute**: Spread data and processing across nodes
 - **Distribute the data**, increasing data capacity
 - **Move processing to the data**, increasing processing capacity
- ◆ **Replicate data**: For availability and processing capacity
- ◆ **Relax Restrictions**: In various ways
 - **Relax** transactional and consistency guarantees
 - **Denormalize** data - sometimes radically
- ◆ **Specialize**: Design for your data and needs
 - Use data models and tools that work well for them

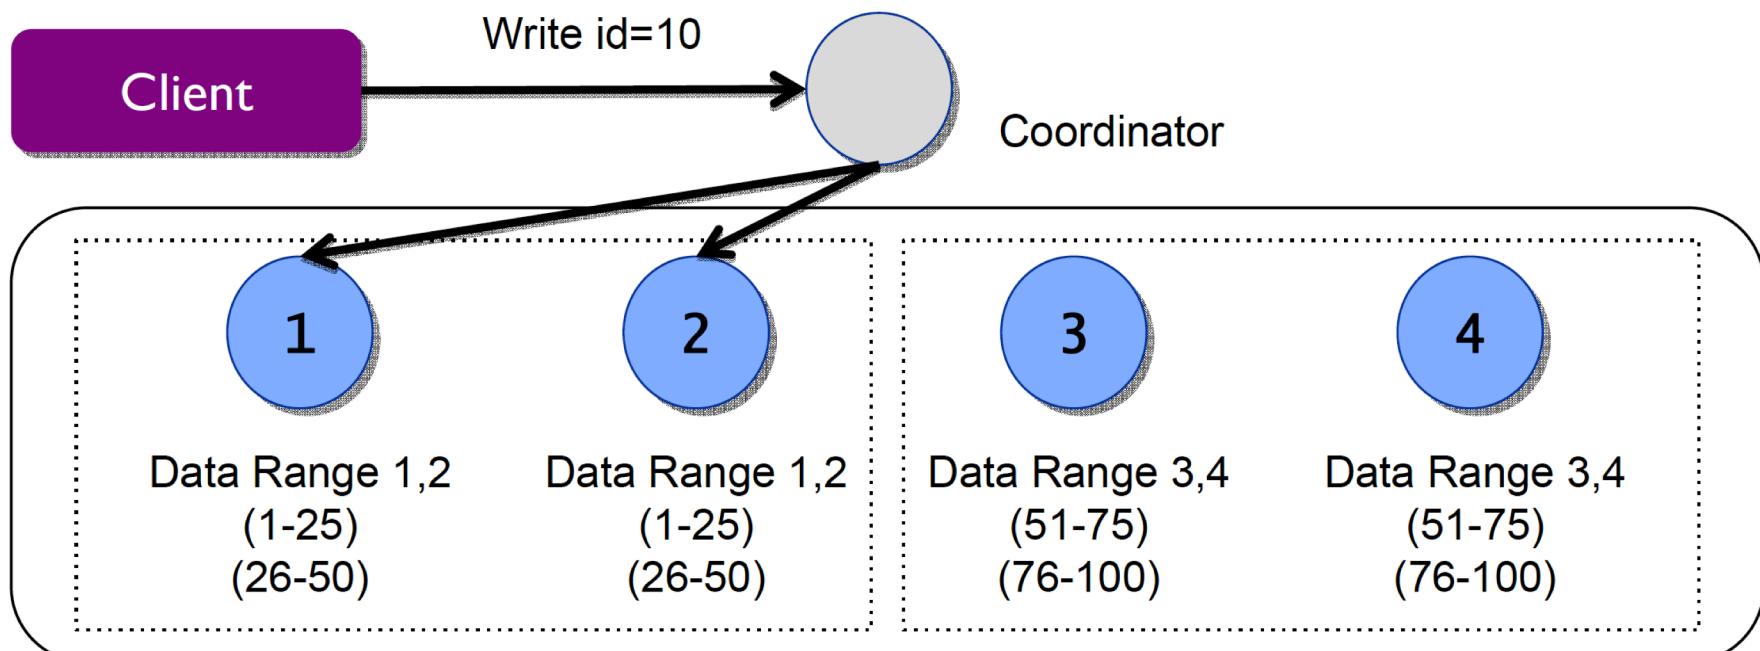
- ◆ **Sharding** partitions large data sets into smaller shards
 - Each shard hold parts of the data, so data is distributed
- ◆ Requires a partitioning scheme to shard data across nodes ⁽¹⁾
 - As well as a way to determine where sharded data goes to
- ◆ Below, 4 nodes each hold $\frac{1}{4}$ of a data range (1-100)
 - An incoming write with id=10 goes to node 1



- ◆ Consider a join that has to first read id=10 and id=70
 - And based on data in the read, reads data with id = 12, 55, and 80
 - Perhaps from another table
- ◆ You end up with many multiple reads across the cluster
 - Lots of network traffic and work to merge all the data into a result
 - Inherent in the relational model, and affects scalability ⁽¹⁾



- ◆ Replication **duplicates data** to multiple nodes ⁽¹⁾
 - Increases scalability - operations can go to all replicas
 - Increases fault tolerance - operations continue if any replica is up
 - Two general types: **master/slave** and **peer-to-peer**
- ◆ Below, a write with id=10 goes to both nodes 1 and 2
 - The data is replicated to both nodes ⁽²⁾



- ◆ **Master/Slave replication:** One node (master) is primary holder of data ⁽¹⁾
 - Updates go to master, and are then propagated to slaves
 - Reads can go to any node (master or slave)
 - Provides good scalability/availability for reads by adding slaves ⁽²⁾
 - Since reads can go to any node with the data

◆ **Advantages:**

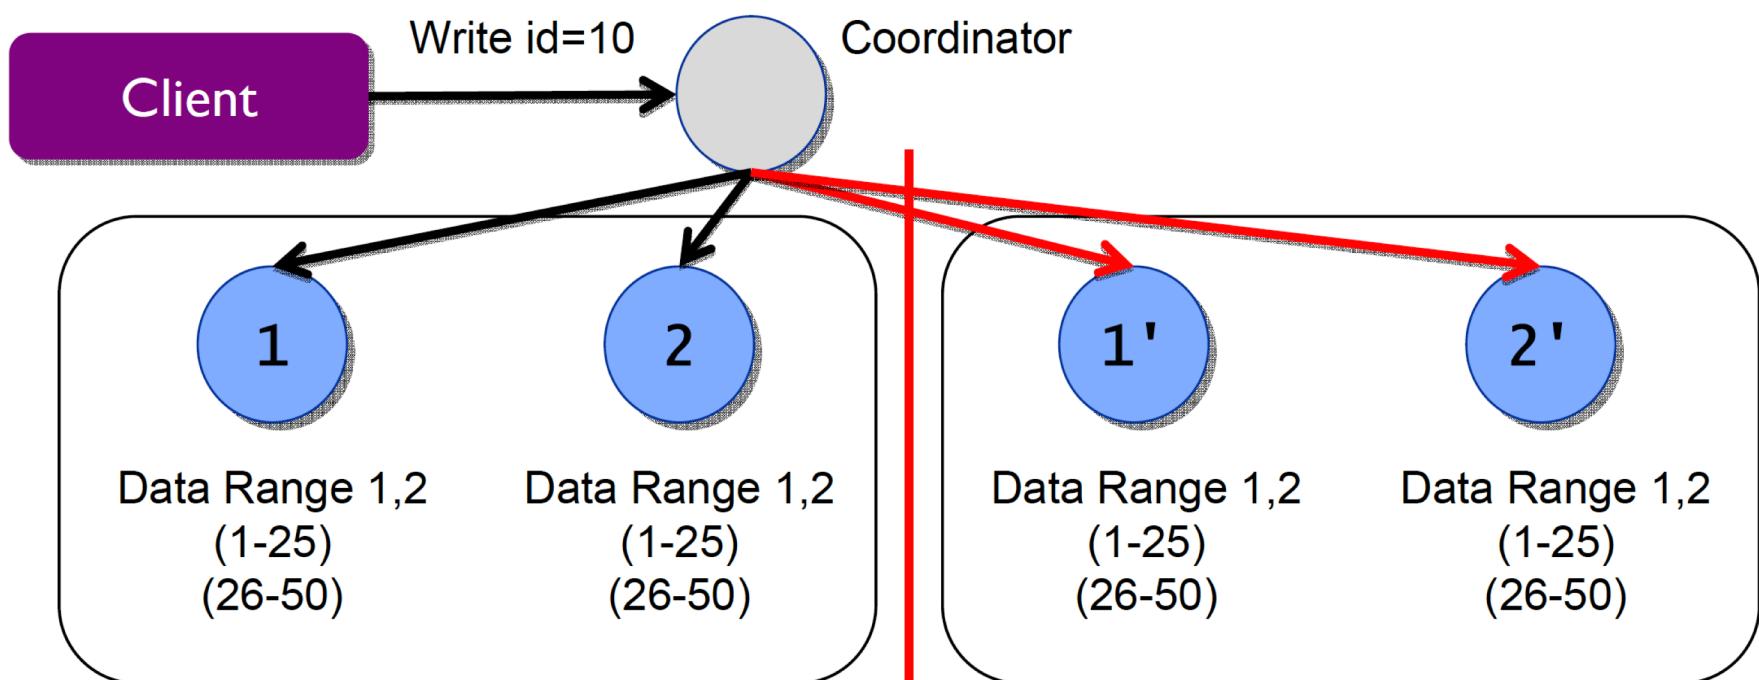
- Reads scale well, and can easily be made highly available
- Can give guarantees on update isolation ⁽³⁾

◆ **Issues:**

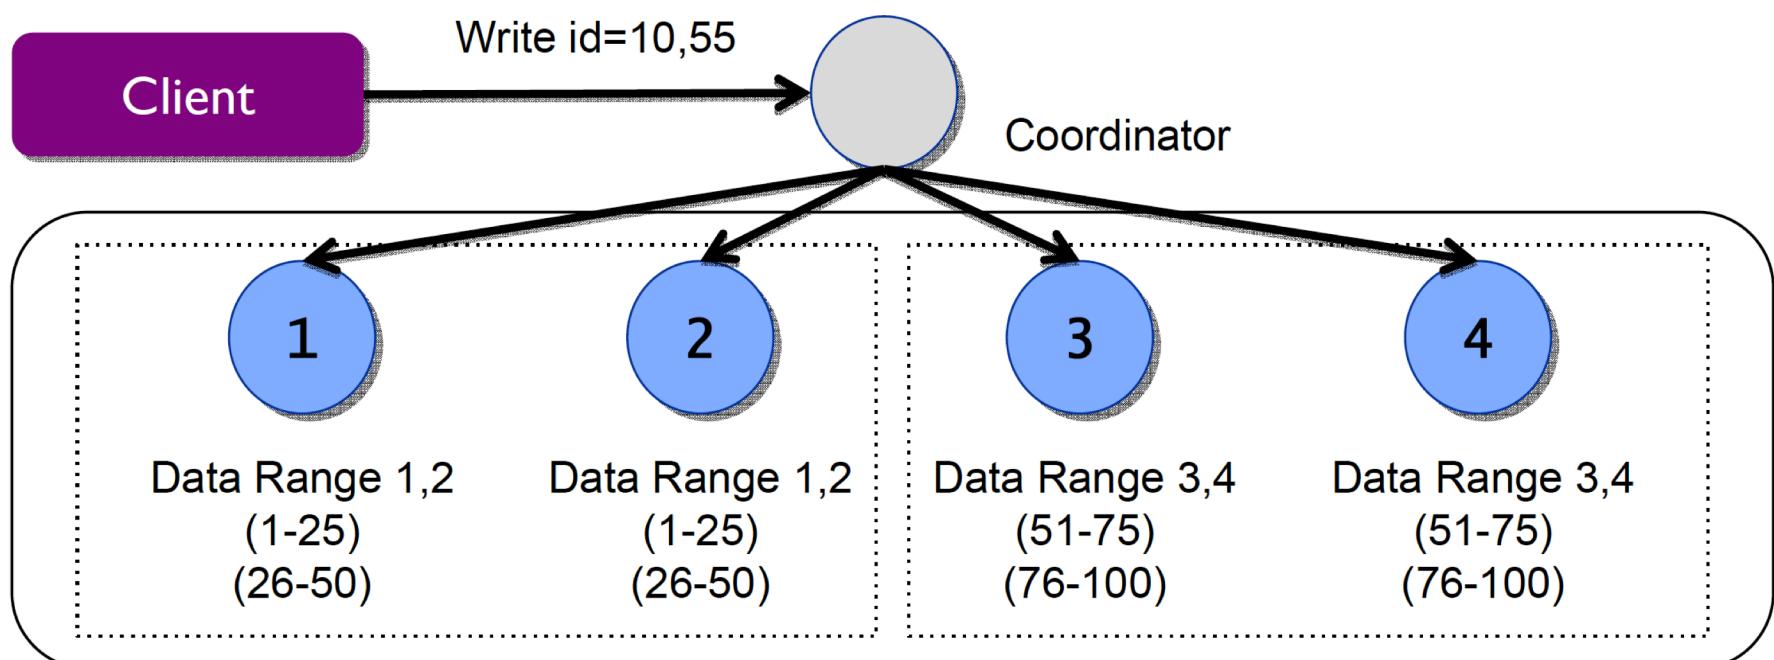
- Updates are limited by master's capabilities, limiting throughput
- If master fails, failover to slave can occur, but may take time or manual intervention, limiting availability

- ◆ **Peer-to-peer replication:** All nodes are equal
 - Reads and writes can go to any node
 - Updates on any node are propagated to all replicas
 - Provides good scalability/availability for reads and writes
 - Both can go to any node with the data
- ◆ **Advantages:**
 - Reads and writes scale well and are highly available
 - Failure of a node doesn't necessarily prevent any reads/writes
- ◆ **Issues:**
 - These solutions are fairly complex
 - Hard to give any guarantees on updates ⁽¹⁾

- ◆ Replication **raises data consistency issues** among replicas
 - e.g. A read occurs before propagation of a write to all replicas ⁽¹⁾
- ◆ Consider two datacenters - with data replicated across them
 - What if the two datacenters can't communicate (a **partition**) ⁽²⁾
 - You now face choices in how you continue to operate (or not)



- ◆ Consider a write that has to write to rows with ids=10 and 55
 - You'll write to all 4 nodes (because of replication and sharding)
 - For ACID guarantees, you'll have to use a complex TX protocol (e.g. 2 phase commit) to insure data integrity
 - The overhead of this protocol will affect scalability, and may add other complexities



Solving The Relational Problem

Compromise: Enhancing Relational to Handle Big Data

- Keep Using Relational Systems
 - Leverage Software Costs / Investment
 - Leverage existing Staff
 - Leverage existing Application Software/Code
- Expand by scaling out (horizontally)
 - Use the Cloud where possible
- Utilize Replication, Sharding
- Leverage Parallelization (queries run in parallel on different nodes)
- Relax ACID compliance where possible, more throughput