

Relational VS NoSQL

Recap

- Relational DBMS software has worked very well for many decades
- Companies have invested lots of money in software built upon relational DBMS infrastructure
- Companies have invested in staff/talent skilled in RDBMS technologies

BUT

- RDBMS systems struggle to adequately scale to support Big Data's volume, variety, and velocity demands
- Big Data is exploding faster than RDBMS technologies can handle

NOSQL Data Models

- **NoSQL** (“Not Only SQL”)
 - Uses clusters:
 - Distribute the Data
 - Distribute the Processing
 - Uses Replication to provide
 - Redundancy
 - High availability
 - Parallel Processing
 - Horizontally scalable

NOSQL Data Models

- **NoSQL** (“Not Only SQL”)
 - Less Structure Required
 - Does not store data in tables with rigid row/column structure
 - Uses Aggregate model
 - Restricts join capabilities
 - Relaxes ACID transaction compliance
 - May use a *non-SQL* query language
 - Typically open source, very low-cost software acquisition

NOSQL Data Models

Relational

- Schema defines rigid structure
 - Tables, Rows, Columns
- Foreign Key relationships
 - Which support joins
- Uses SQL
- Maintains ACID compliance
- Normalized: store a value only once
- Clustering available, but challenging
- Leading DBMS solutions are quite expensive

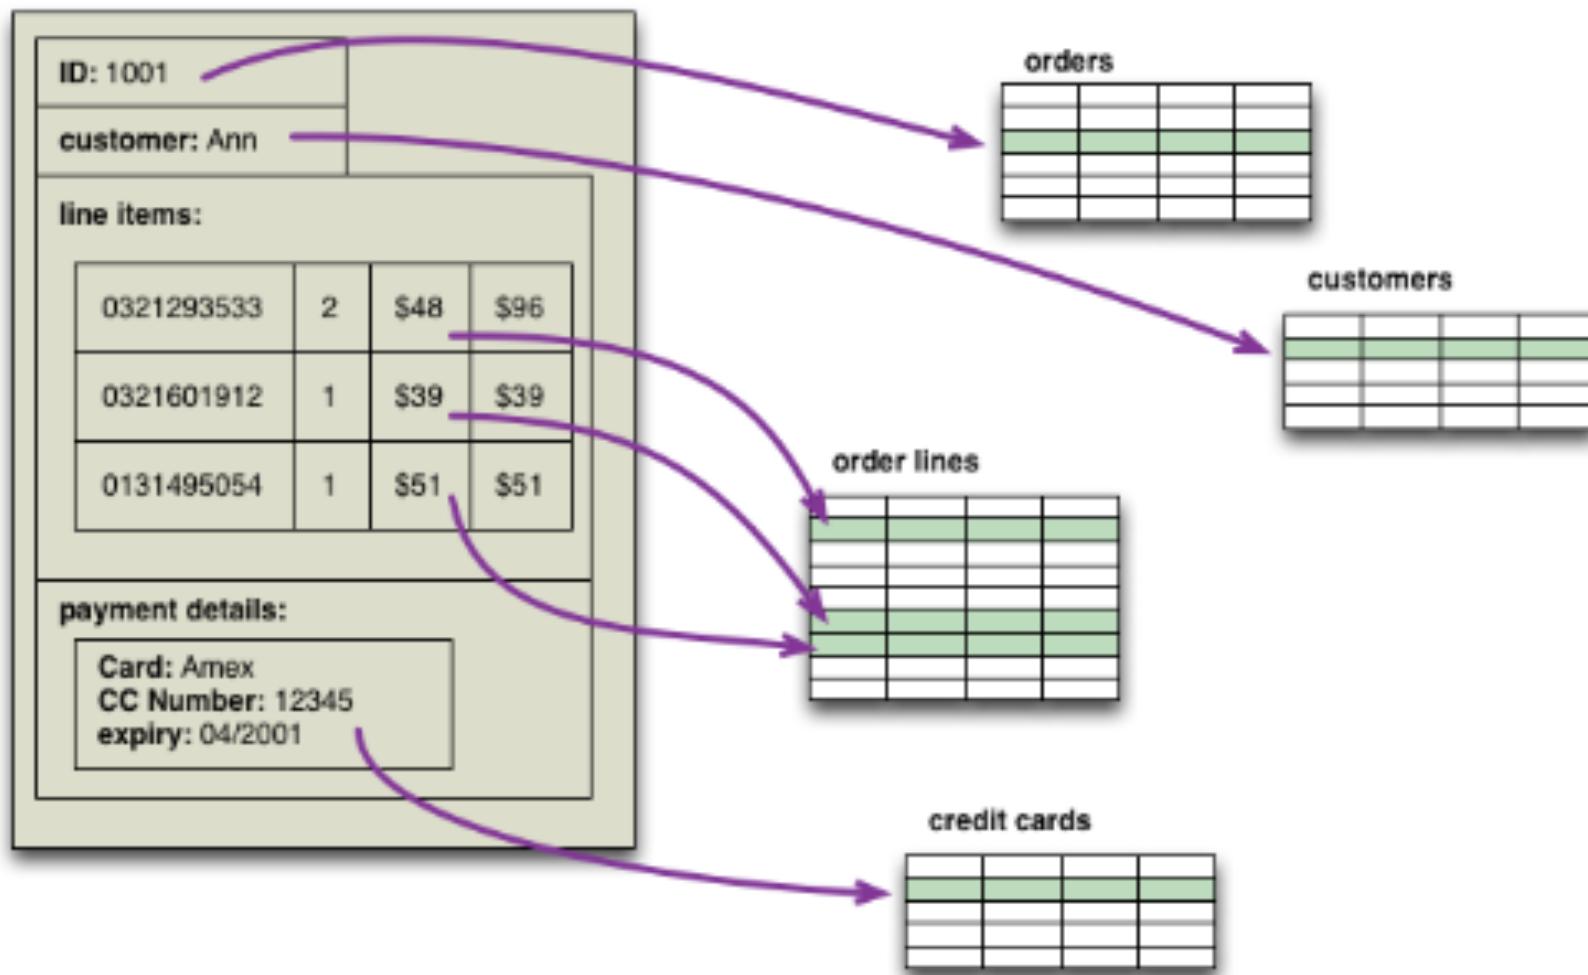
NoSQL

- Stores related values in aggregates
- Flexible structure:
 - Ranges from none to some
- No joins
- Relaxed ACID compliance
- De-normalized
- Uses alternate query language
- Designed to support clustering
- Almost all players are open source and low-cost

NOSQL Data Models

- **NoSQL “aggregate” Data Model**
 - RDBMS requires Tables, Rows, Columns as data stores
 - Columns have “domains”
 - Third normal form – no multi-values, “store it once”
 - NoSQL systems use AGGREGATES as data stores
 - Data values are grouped together as users need them
 - Think of an unnormalized document, or an “object”
 - Contains related values that are retrieved and manipulated together

NOSQL Data Models



- **Aggregates are conceptually the opposite of 3rd Normal Form**
- **Why aggregates?**
 - It is difficult to spread a relational model across nodes in a cluster
 - Replication and sharding introduce big challenges in consistency
 - Each query should minimize the number of nodes being accessed across the cluster
 - Data values that are accessed together should live on the same node

NOSQL Data Models

- **NoSQL Data Schema/Models**

- Document Store (using XML or JSON format)
- Graph (using node/edge structures with properties)
- Key-Value pairs
- Wide Columnar store (rows with dynamic columns holding key-value pairs)

- Document
 - MongoDB
- Graph
 - Neo4j
- Key-Value
 - Amazon Dynamo
- Wide Column
 - Google BigTable, Apache Cassandra

NOSQL Data Models

- **Document Database**
 - Organized around a “document” containing text
 - Can handle very large data volumes
 - Provides Speed and Scalability
 - Document format is easily understood by humans
 - No “schema”, but JSON/XML provides internal structure within a document
 - Documents are indexed and stored within “collections”
 - Supports full text search
- **Popular Implementations** (open source)
 - MongoDB
 - CouchDB

JSON

```
{  
  _id: ObjectId(7df78ad8902c)  
  title: 'MongoDB Overview',  
  description: 'MongoDB is no sql database',  
  by: 'tutorials point',  
  url: 'http://www.tutorialspoint.com',  
  tags: ['mongodb', 'database', 'NoSQL'],  
  likes: 100,  
  comments: [  
    {  
      user:'user1',  
      message: 'My first comment',  
      dateCreated: new Date(2011,1,20,2,15),  
      like: 0  
    },  
    {  
      user:'user2',  
      message: 'My second comments',  
      dateCreated: new Date(2011,1,25,7,45),  
      like: 5  
    }  
  ]  
}
```

XML

```
<contact>
    <firstname>Bob</firstname>
    <lastname>Smith</lastname>
    <phone type="Cell">(123) 555-0178</phone>
    <phone type="Work">(890) 555-0133</phone>
    <address>
        <type>Home</type>
        <street>123 Black St.</street>
        <city>Big Rock</city>
        <state>AR</state>
        <zip>23225</zip>
        <country>USA</country>
    </address>
</contact>
```

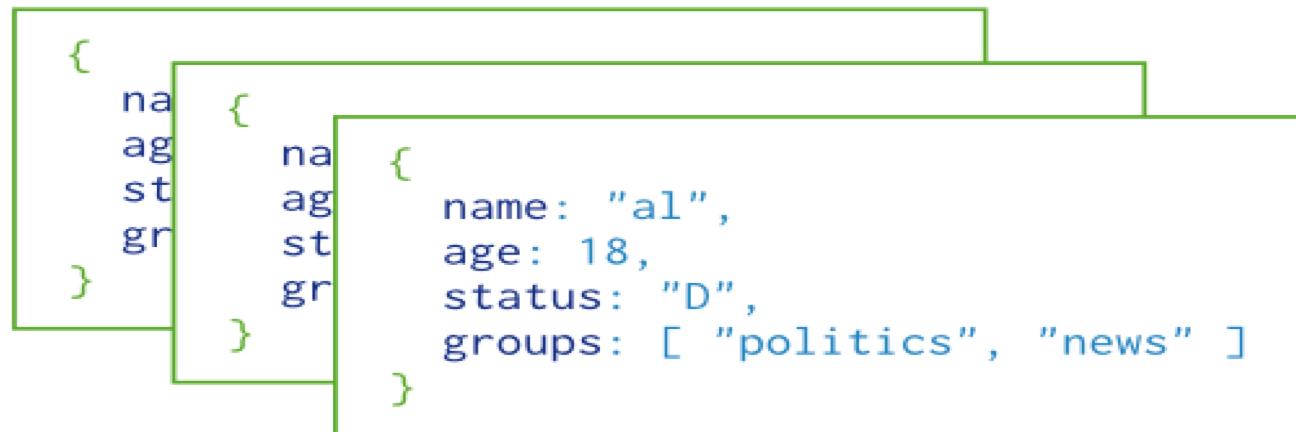
- **Document Database**

- Keeps related information together (not normalized into tables)
- Access to a document is fast (index/key/URL)
- ACID compliance is maintained only within a document
- Cannot easily “join” across documents
- Documents are kept in “collections”

NOSQL Data Models

```
{  
  name: "sue",  
  age: 26,  
  status: "A",  
  groups: [ "news", "sports" ]  
}  
← field: value  
← field: value  
← field: value  
← field: value
```

A MongoDB document.



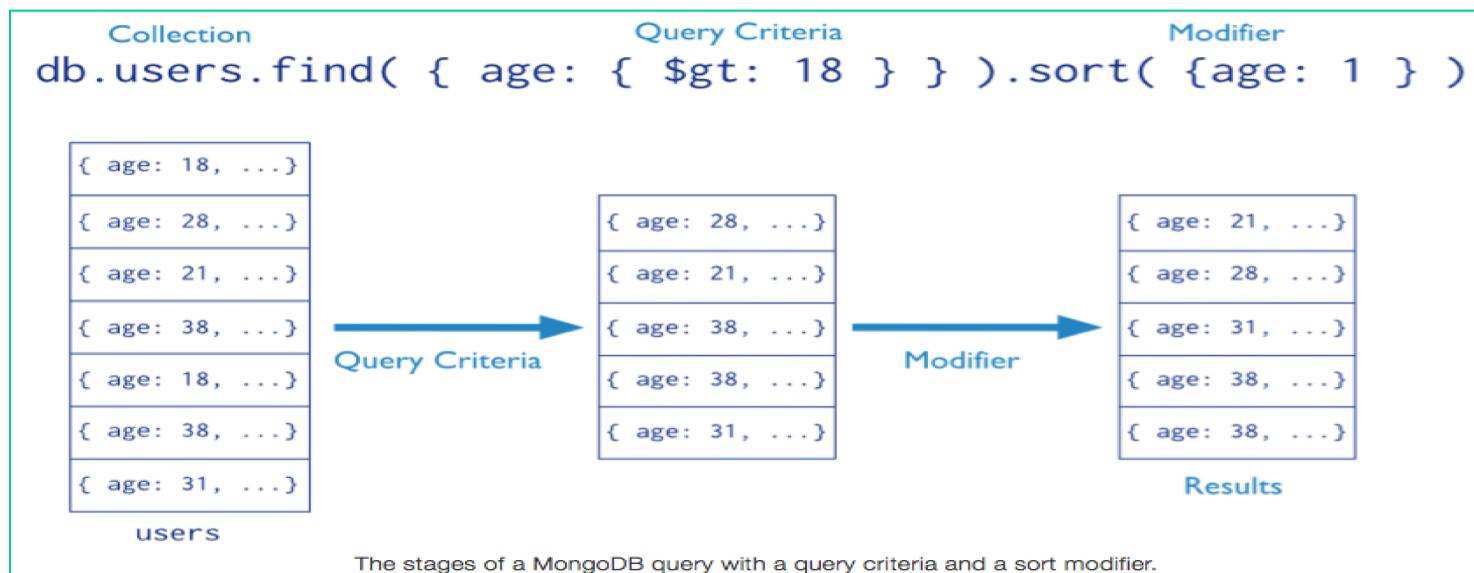
Collection

A collection of MongoDB documents.

NOSQL Data Models

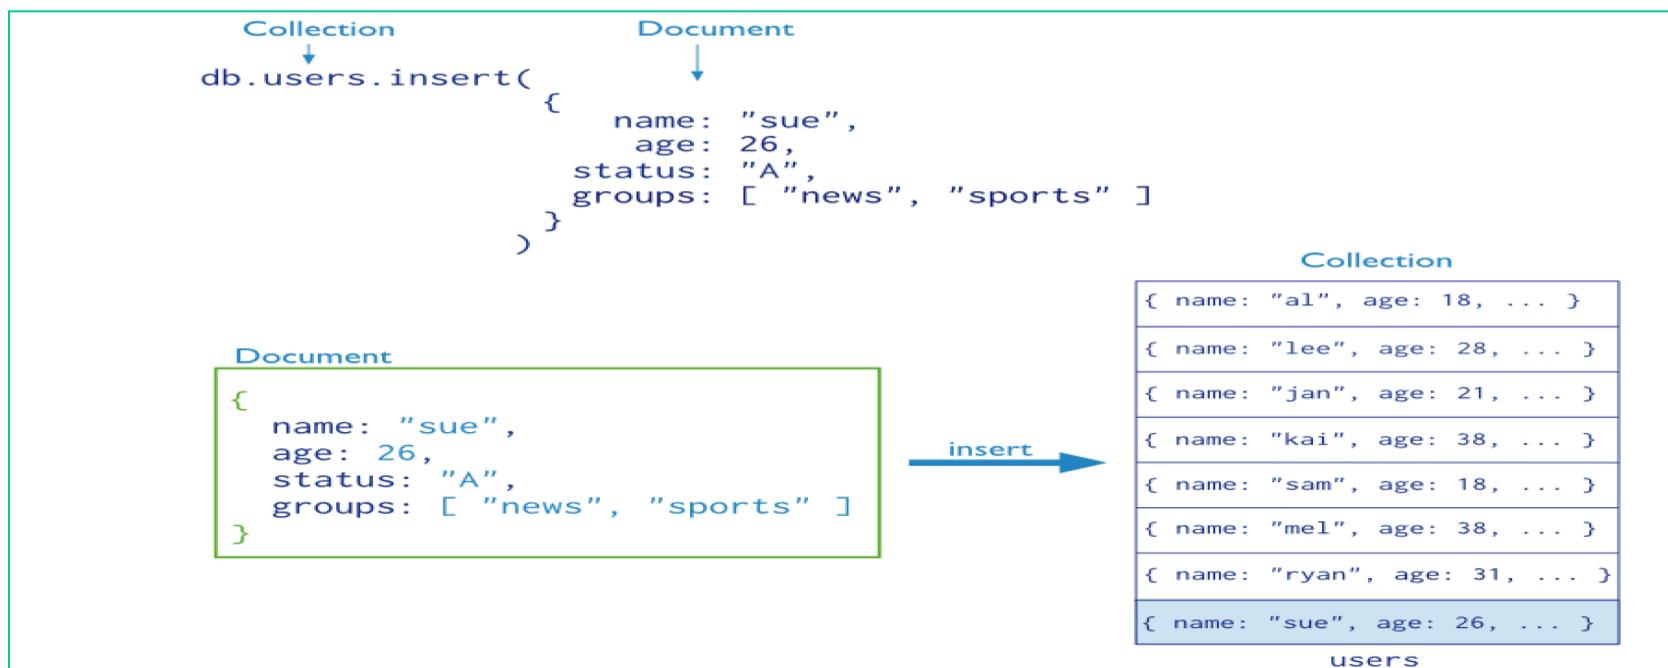
In MongoDB a query targets a specific collection of documents. Queries specify criteria, or conditions, that identify the documents that MongoDB returns to the clients.

A query may include a *projection* that specifies the fields from the matching documents to return. You can optionally modify queries to impose limits, skips, and sort orders.



NOSQL Data Models

Data modification refers to operations that create, update, or delete data. In MongoDB, these operations modify the data of a single collection. For the update and delete operations, you can specify the criteria to select the documents to update or remove.

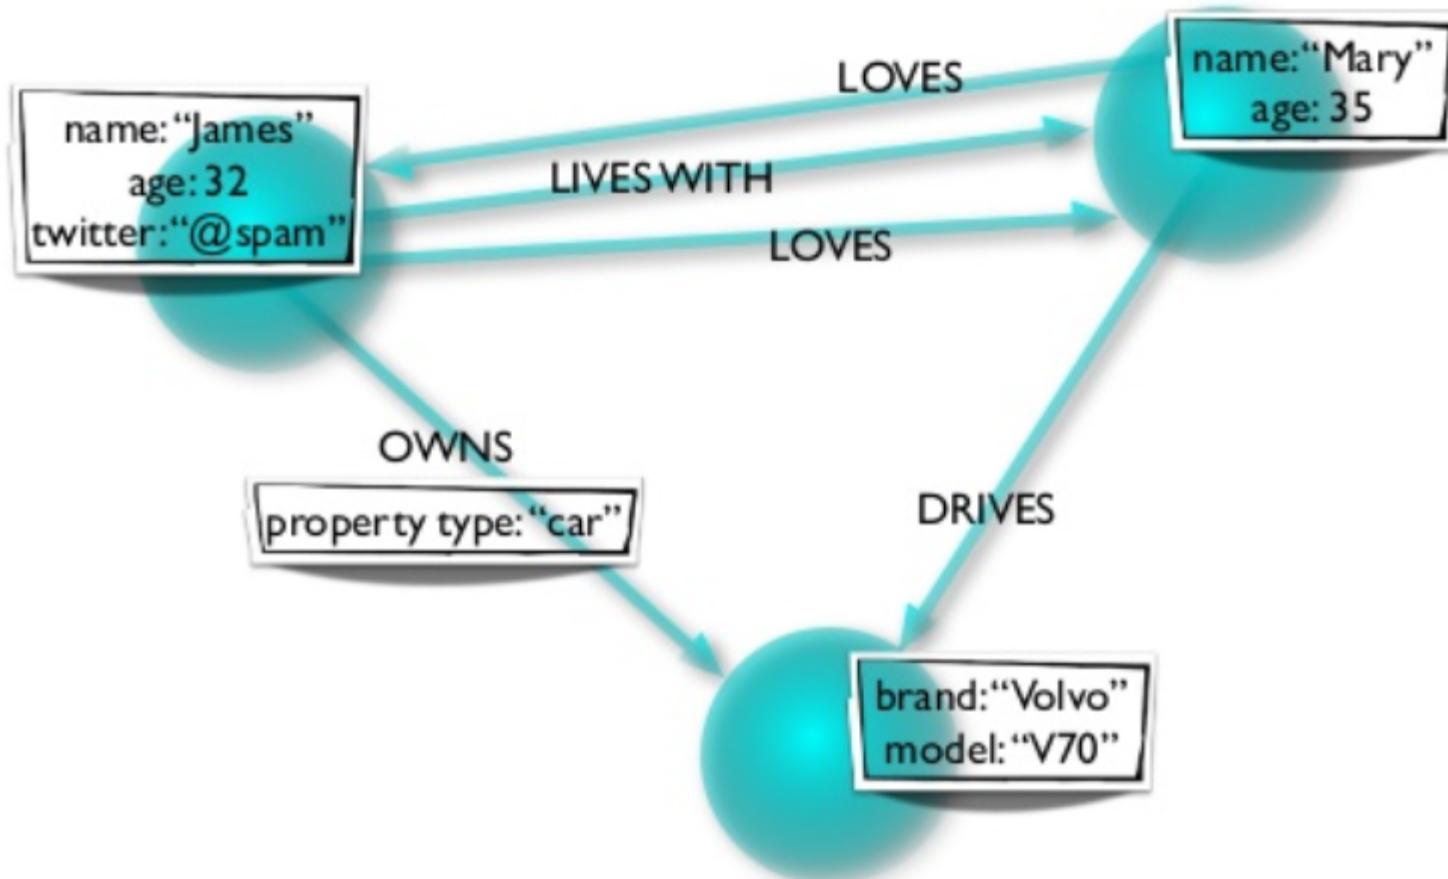


NOSQL Data Models

- More on MongoDB next week

- **Graph Database**
 - Uses a graph structure consisting of
 - Nodes – Represents an entity (like a person)
 - Edges – Represents a relationship between entities
 - Properties – Attributes associated with Nodes and Edges
 - Supports navigation along edges from a starting point node
 - Designed for applications tracking the inter-connections among entities
 - Who is friends with whom? (In a social network application)
 - Who is following me, who am I following?
 - Uses a pattern matching query language to navigate nodes & edges
- **Popular Implementations** (open source)
 - Neo4j

NOSQL Data Models



NOSQL Data Models

- **Stopped here on Friday April 3**

- **Key-Value Pairs Database**
 - “Schemaless”
 - Maps a key to an opaque value
 - That is, the database doesn’t understand anything within the value
 - Simple operations (put, get, remove, modify)
 - Keys are unique in a collection
 - Most useful when access is by the primary key
 - May be a building block for other data models (such as key-value pairs within a wide-columnar store like Cassandra)
- **Popular Implementations**
 - Amazon Dynamo (available via AWS in the cloud)
 - Riak, Redis (open source)

Product Catalog Example

```
{  
    Id: 206,  
    Title: "20-Bicycle 206",  
    Description: "206 description",  
    BicycleType: "Hybrid",  
    Brand: "Brand-Company C",  
    Price: 500,  
    Color: ["Red", "Black"],  
    ProductCategory: "Bike",  
    InStock: true,  
    QuantityOnHand: null,  
    RelatedItems: [  
        341,  
        472,  
        649  
    ],  
    Pictures: {  
        FrontView: "http://example.com/products/206_front.jpg",  
        RearView: "http://example.com/products/206_rear.jpg",  
        SideView: "http://example.com/products/206_left_side.jpg"  
    },  
    ProductReviews: {  
        FiveStar: [  
            "Excellent! Can't recommend it highly enough! Buy it!",  
            "Do yourself a favor and buy this."  
        ],  
        OneStar: [  
            "Terrible product! Do not buy this."  
        ]  
    }  
}
```

Key-Value Pairs Database Example (Amazon DynamoDB)

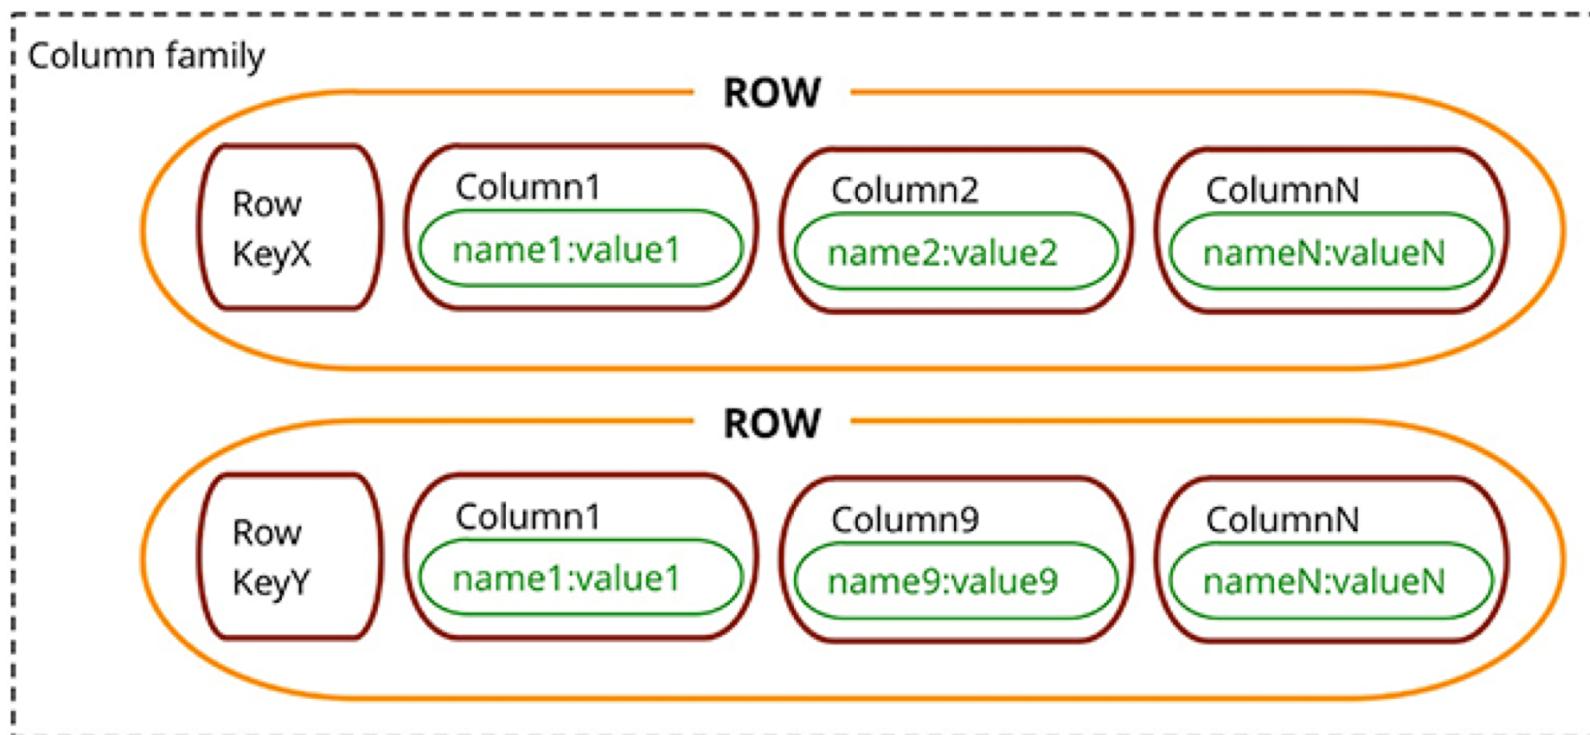
- The key value (Id) is 206.
- Most of the attributes have simple data types, such as String, Number, Boolean and Null.
- One attribute (Color) is a String Set.
- The following attributes are document data types:
 - A List of Related Items. Each element is an Id for a related product.
 - A Map of Pictures. Each element is a short description of a picture, along with a URL for the corresponding image file.
 - A Map of ProductReviews. Each element represents a rating and a list of reviews corresponding to that rating. Initially, this map will be populated with five-star and one-star reviews.

- **Wide-Column (Column Family) Store Database**
 - A TWO-LEVEL aggregate
 - Data is stored within “collections” of dynamic related columns
 - Similar to key/value with the pairs having columnar structure
- **Based on Google’s “Big Table”**
- **Popular Implementations** (open source)
 - Cassandra
 - HBase

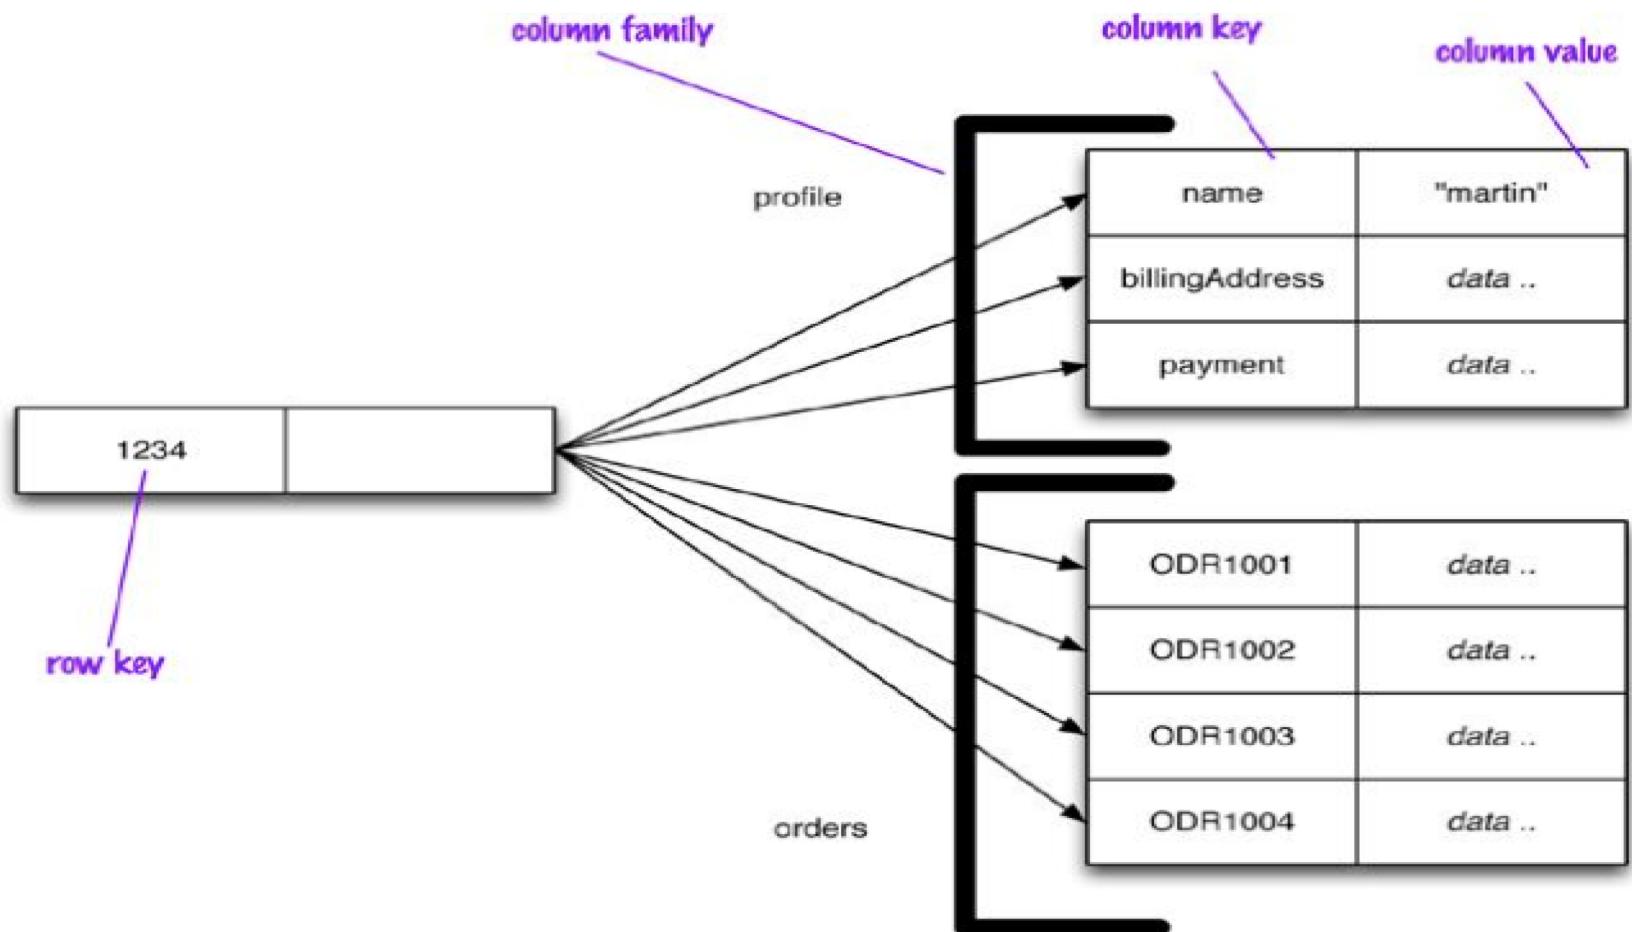
NOSQL Data Models

- **The first key is the row-key**
 - The entire row is an aggregate of related data
 - The row consists of many key-value pairs (“columns”)
- **The second key is the column family**
 - Each column family consists of sparse key-value pairs
 - “Sparse” means the column value isn’t stored if it isn’t needed

NOSQL Data Models



NOSQL Data Models



- **Wide-Column Store Example**
- Consider an application that needs to store stock trades and retrieve them by stock symbol
- Some data that might be stored includes the following:
 - symb: The stock symbol
 - id: A unique trade id
 - date: Date of the trade
 - price: Trade price
 - quant: Number of shares traded
 - notes: General comment field.

NOSQL Data Models

- Below, we illustrate trade data in a wide column store
 - The stock symbol (symb) is the **row key**
 - Generally the complete row is stored on a **single node**
 - The trades are stored as **columns** in the row
 - If there were many trades, the row could get very wide
 - Note also, that Trade 2 has no notes column
 - Contrast this to a normalized relational model with a trade table which would generally store trades in separate rows

Trade 1						Trade 2				Trade 3			
symb	id	date	price	quant	notes	id	...	quant	id	...	notes		
GM	1c3f	2014-09-15	32.96	100	blah	2a5f		400	2e99		bar		

Benefits of a Wide Column Store

- Lookup of rows by row key can be **very fast**
- In a distributed cluster system, the complete row is stored on one node
- May be stored redundantly across the cluster
- Can be retrieved with one access
- Good fit for **scale-out** systems
- Can scale to large capacity and high availability
- The columns are "**sparse**"
 - That is, columns with no values are absent, saving space
 - As shown in Trade 2 having no Notes column
- Flexible access to column data in a row
 - You can flexibly query on them
 - e.g. - get all trades, get the last 10 trades, or 1st ten trades, etc.

Issues with a Wide Column Store

- The data model is **complex**, and **not very intuitive**
- Generally, you create your data model **based on your queries**
 - Possibly taking into account how data is distributed
- For example, querying for trades by stock symbol is very fast for the trade table shown earlier -- it's just one query to one node
- Consider finding all trades (for all stocks) on a particular date
 - The previous data model is not very good for this
 - You'd need to query every node, get the trades from that node, then combine them all for the result
- In practice, that may not scale
- You could, however, store data optimized for retrieval by date
- This would require another table / column layout

Issues with a Wide Column Store

- Client code **may be extensive**
 - For example, to keep track of trades by both stock symbol and trade date, you may need to maintain two sets of data
 - Which needs to be done via application code on a client
- Different from relational model where you maintain one set of data, and just write queries as needed
- However, wide column stores tend to have much more scale out capability -- Which is why we use them in the first place

Linear Scalability as in Cassandra

- Capacity may be easily added simply by adding new nodes
- For example,
 - Suppose 2 nodes can handle 100,000 transactions per second,
 - 4 nodes will support 200,000 transactions/sec, and
 - 8 nodes will tackle 400,000 transactions/sec:

