

MongoDB is a NoSQL “Document” database.

- Stores collections of documents in a key:value pair format
- MongoDB is NOT Relational
- MongoDB does not store data in tables
- MongoDB does not use the SQL query language
- MongoDB uses a JS-like query language
- Community edition is free, Enterprise edition = \$\$\$
 - Enterprise users can purchase support, advanced features, cloud deployment

MongoDB Concepts

- A MongoDB **DATABASE** can contain one or more **COLLECTIONS**
- A MongoDB **COLLECTION** can contain many **DOCUMENTS**
 - Each **document** has a primary key
 - Primary keys (and any other field) can be **indexed** for faster performance

MongoDB Concepts

- Documents are structured and stored in JSON-like format
 - MongoDB stores data in BSON
 - Binary JSON
 - BSON was invented by the MongoDB folks...

JSON

- Invented along with Java Script
- Pushed out and replaced XML
 - XML is more verbose
 - JSON is easier for humans to work with
- A string “key” is mapped to a “value”
 - The value can be a number, string, array
- Very widely used in web-based software development

BSON

- JSON is string based
- String parsing is relatively slow
- JSON is human-readable, but less efficient for storage and movement of data across networks
- BSON was invented by MongoDB folks to replace JSON with a data format/structure that provides
 - Faster data movement
 - More efficient storage
 - More flexible – can store more types of data (like integer versus floating point)
- BSON encodes data item TYPE and LENGTH in binary notation

Notes on MongoDB

```

{"hello": "world"} → \x16\x00\x00\x00      // total document size
                    \x02          // 0x02 = type String
                    hello\x00     // field name
                    \x06\x00\x00\x00world\x00 // field value
                    \x00          // 0x00 = type E00 ('end of object')
```

```

{"BSON": ["awesome", 5.05, 1986]} → \x31\x00\x00\x00
                                   \x04BSON\x00
                                   \x26\x00\x00\x00
                                   \x02\x30\x00\x08\x00\x00\x00awesome\x00
                                   \x01\x31\x00\x33\x33\x33\x33\x33\x33\x14\x40
                                   \x10\x32\x00\xc2\x07\x00\x00
                                   \x00
                                   \x00
```

MongoDB Documents

- They are “polymorphic” (i.e. “multiple shapes”...)
 - i.e. Not all documents in a collection must have all the same key:value pairs.
- There is no need to declare the structure of a document
 - It is “self-describing”
 - The database takes the JSON you provide, converts it to BSON, and stores/retrieves it efficiently
 - Converts back to JSON on retrieval
- You can add fields to one document in a collection without modifying any other documents in the same collection

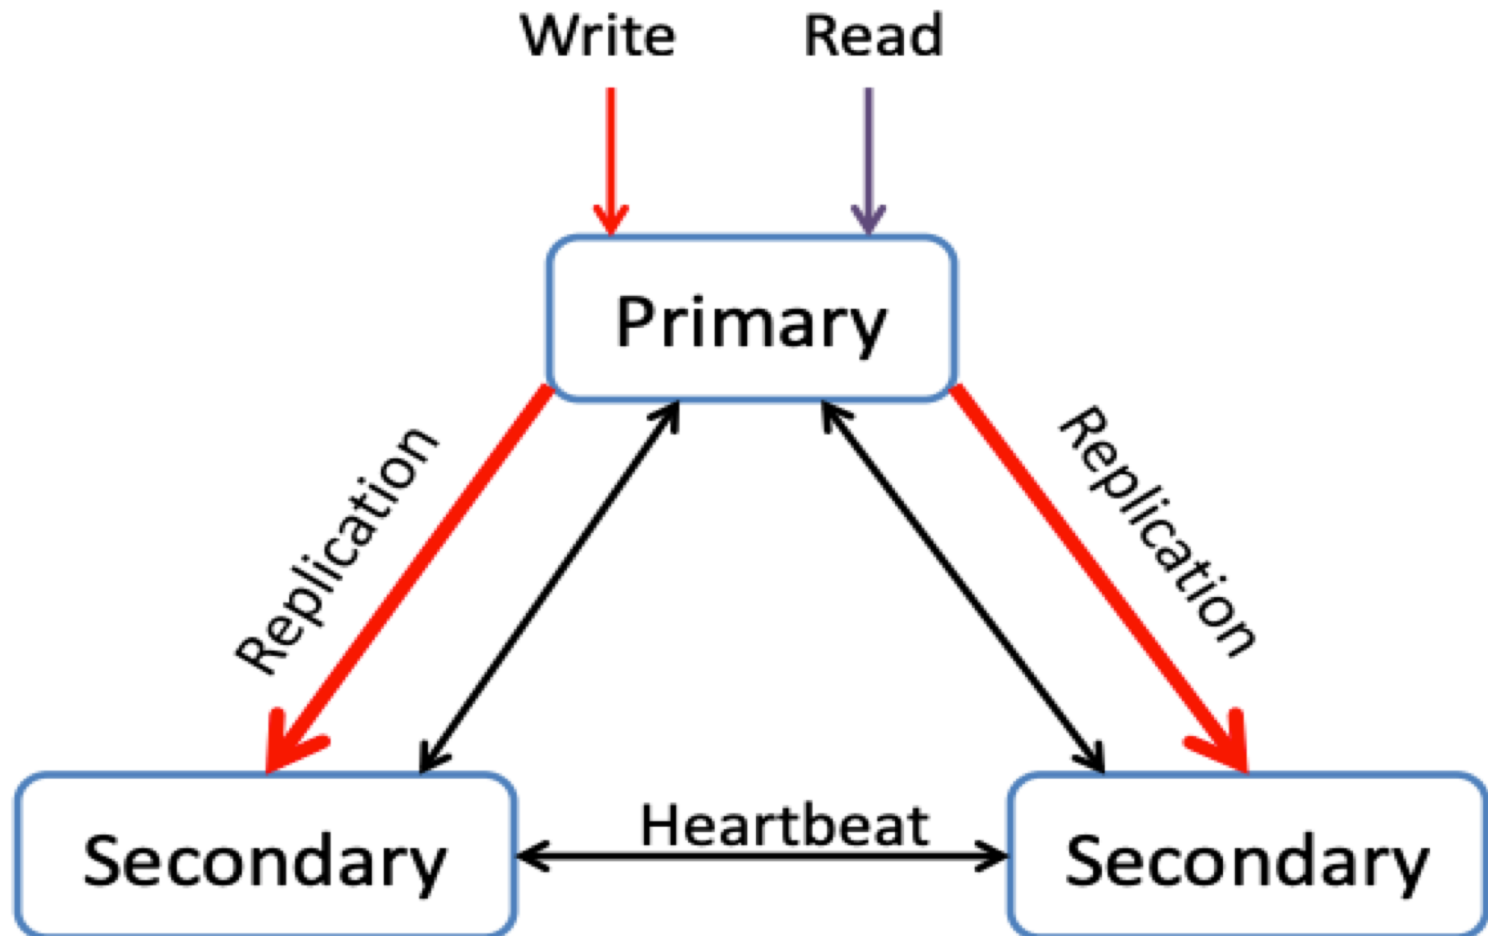
MongoDB relies on REPLICATION

- MongoDB provides horizontal scaling
 - You can configure a scalable number of nodes in a cluster
 - The cluster can be spread across data centers and geography
 - MongoDB can easily scale READ operations across the cluster (parallelization)

MongoDB relies on REPLICATION

- You can configure a number of replica sets
 - Each set is replicated across the cluster
 - Provides High Availability
 - If MongoDB detects that it has lost a node, it will shift processing over to another replica

Replication



MongoDB node role management

- If the Primary node fails for any reason, the other members vote to elect a new primary from among the secondary nodes

No Downtime for Upgrades

Distribution via Replication allows administrators to

- take a node offline
- upgrade hardware or software
- bring it back online
- NO DOWNTIME needed

MongoDB relies on SHARDING

- MongoDB uses horizontal scaling for sharding
 - You can spread your data out across multiple nodes
 - A shard is referred to as a “partition”
 - Distributes WRITES across multiple nodes/partitions (parallelization)
 - Data can be distributed based on user query patterns


MongoDB relies on SHARDING

- MongoDB offers architects options regarding sharding
- Each document has a primary key
 - Partition by key ranges
 - Co-locating documents based on geography
 - Partition by hash
 - A more random but uniform spread of data

Partition

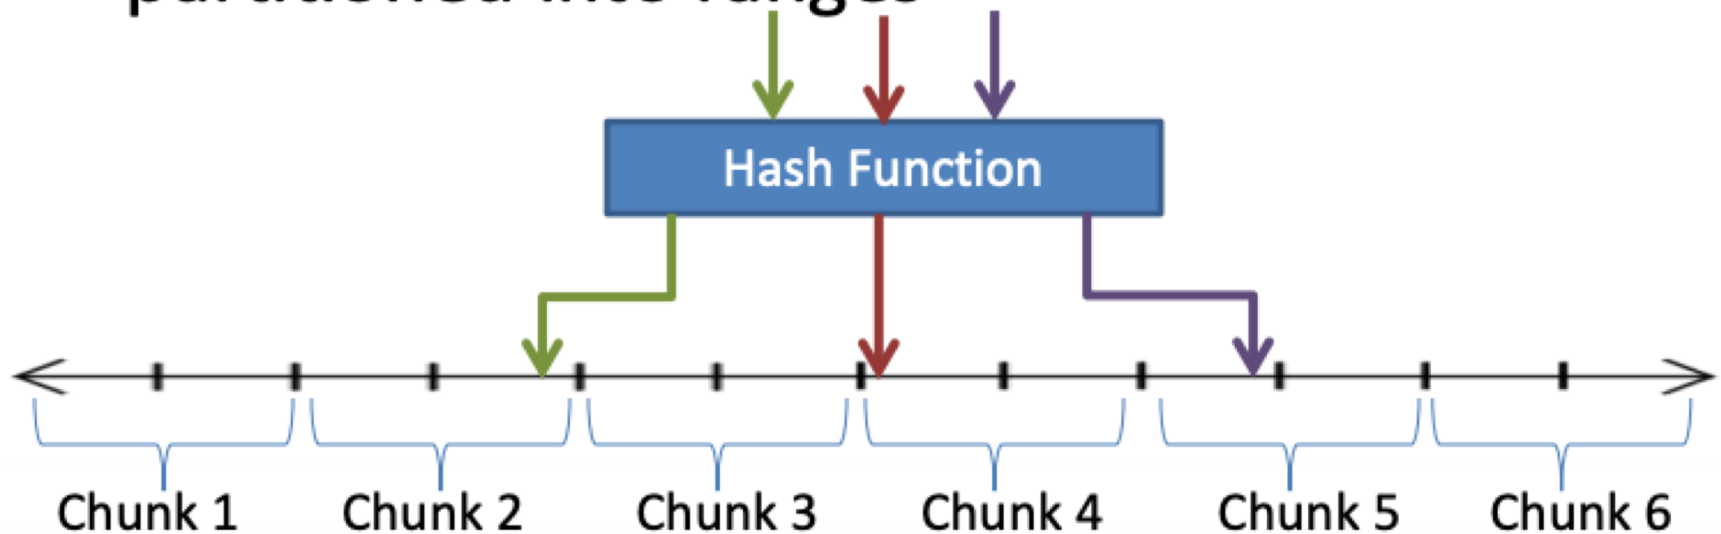
- Shard Key: Single or compound field in schema used for data partitioning
- Partitions are called *chunks*. Two strategies:
 - Range based: Shard Key Values are partitioned into ranges

Total Key Space for x



Partition

- Hash based: Hash of shard key values are partitioned into ranges



- Hash Scheme leads to better data balancing

Balancing

- Splitting: Background process which splits when a chunks grows beyond a threshold
- Balancing: Migrates chunks among shards if there is an uneven distribution

More Resources

- <http://mongodb.org>
- <https://docs.mongodb.com/manual/tutorial/getting-started/>